

## 5 Exception safety

### Handling errors and exceptions

UNIVERSITY OF HELSINKI

1

## Preview

- different levels of exception safety
- support mechanisms for exception safety
- implementing exception safety
- exception safety in the standard C++ library

UNIVERSITY OF HELSINKI

2

## Background

- originally, exceptions and error handling were very poorly understood aspect of C++
  - exception safety issues and requirements were seen and added to the C++ standard at very last moments
- Stroustrup's own C++ book discusses exception safety in an extra Appendix E: *Standard-Library Exception Safety* - (also available from <http://www.research.att.com/~bs/>)
- two concerns:
  - class invariants must be maintained / restored
  - no resources may be leaked
    - including: memory space, opened files, locks, connections, and any other system resources

UNIVERSITY OF HELSINKI

3

## Background (cont.)

- exception safety is essential for reusable units (libraries)
- the C++ standard library provides at least the basic exception guarantee: invariants are maintained, and no resources are leaked
- of course, the same should hold for all reusable library components (own or third-party)

### Note

- C++ standard libraries still do not necessarily check for all errors (following C-style error handling strategy)
- but when they check - and potentially throw exceptions - they *do* guarantee exception safety
- for many simple (say, only a single added element) operations, the C++ standard library guarantees *strong exception safety* (discussed later)

UNIVERSITY OF HELSINKI

4

## Revision on invariants

- use defensive programming and self-checking objects
- a *class invariant* is an assertion that holds before and after any operation manipulating an object
- preconditions tests external failures which the unit cannot handle itself: must throw an exception

```
if (!precondition || other external failure)
    throw AnException("diagnostics"); // to the caller
```
- invariants and (concrete) postconditions test internal states that don't make sense to outsiders, and may indicate a bug in code => use asserts to eliminate them

```
assert(isInvalidInternalState_); // aborts if not
```
- often, we don't know the original reason of a failure: perhaps a programming error or some external factor

UNIVERSITY OF HELSINKI

5

## Revision on invariants (cont.)

- preconditions and external failures provide a pragmatic trade-off what to check, at the boundary of a unit
- note that the C++ standard library uses the same strategy (checks for selected operations and may throw)

### The *fundamental problem* with exception safety

- an exception thrown from some component or function may interrupt the algorithm and leave the state of the calculation (objects) in some indeterminate state
  - the class invariant does not hold => the object cannot be even *destructed* without causing undefined behavior

UNIVERSITY OF HELSINKI

6

## Many sources of exceptions

- user-supplied and system functions, such as allocator functions, can throw exceptions (from [Stroustrup])

```
void fun (std::vector<X>& v, X const& x) {  
    v [2] = x;           // X's assignment may throw  
    v.push_back (x); // vector <X>'s allocator may throw  
    std::sort (v.begin (), v.end ()); // less-than may throw  
    std::vector<X> u = v; // X's copy ctor may throw  
    ...  
} // u is destructed here: X's dtor should not throw!
```



7

## Levels of exception safety

1. **Basic Guarantee**: no leaks, and maintains invariants.
2. **Strong Guarantee**: succeeds, or leaves state unchanged.
3. **Nofail Guarantee**: doesn't fail in any circumstances.

- the last one (*Nofail*) is often needed to implement the former ones; e.g., an assignment of a primitive value (pointer) cannot fail
- the *strong guarantee* for a complicated update may require a "roll-back" mechanism (can be too expensive)
- "*maintaining invariants*" means
  - the object is in *some* valid state (but not necessarily in the one we would like it to be)
  - but it can be at least released and *destroyed*



8

## Exception safety (cont.)

- e.g., `std::vector<T>::push_back` is designed to give the *strong guarantee*: the item is added or no change

Additionally:

4. **Exception Neutrality**: exceptions originating from components are always passed through unmodified
  - relevant for a container handling and copying its elements, especially when using C++ templates
  - e.g., standard `vector<T>::push_back` also manifests exception neutrality: after any internal clean-up, propagates the original exception caused by the copying operation - that depends on the actual element type (T)



9

## Notes on exception safety levels

- exceptions are necessary for making reusable libraries and components work: a component may detect an error (violation of invariant) but doesn't know how to handle it
- exception safety means the capability to handle a throw caused by a failure, especially to manage resources

### No exception safety

- a failed operation may leave an object in an invalid state (breaking class invariant) and/or leak resources
- this strategy may well be OK: just exit and let the user run the program again

### Basic guarantee

- maintain class invariants and don't leak resources
- the processing might possibly be resumed; *at least* the object can be destroyed when propagating the exception



10

## Levels of exception safety (cont.)

### Strong guarantee

- either an operation succeeds or it doesn't cause any changes (IO operations should behave in a similar way: either succeed, or leave the variable untouched)

### Nofail guarantee (also called: "Nothrow guarantee")

- an operation cannot ever fail (and throw); e.g., assignment of a primitive value (say, a pointer) cannot fail; the STL container functions `size` or `swap` do not throw
- especially, destructors must provide the nofail guarantee

The basic strategy for exception safety

- first calculate results separately; this may succeed or fail (and throw)
- if calculation succeeds, only then make changes in a safe way (that cannot cause any throw)



11

## Calling constructors and destructor

- constructors & destructors are usually called by compiler  
`X * ptr = new X; // reserve memory, then construct X`  
`delete ptr; // destruct X, then release memory`
  - note that if `p` is zero (0), `delete` has no effect
- when necessary, *allocation* can be separated from object *initialization* with the so-called *placement-new* operator  
`void * p = ::operator new (sizeof(X)); // allocate space`  
`ptr = new (p) X; // placement-new constructs X at p`
- similarly, we can separate destruction & deallocation  
`ptr->~X (); // destruct the object pointed by ptr`  
`::operator delete (ptr); // delete operator frees space`
  - since destructor is a member function, you can call it explicitly - but then must ensure that compiler doesn't!



12

### Language support for exception safety

- C++ language rules ensure that exceptions thrown while construction will be handled correctly
  - either the object is *fully built* (its invariants OK), or its members become (automatically) destructed
- also *new* operations are implemented safely; "p= new T;" is compiled into something like:

```
p = ::operator new (sizeof (T)); // may fail & throw
try {
    new (p) T; // placement-new: create a T here
} catch (...) { // note the exception neutrality
    ::operator delete (p); // release raw memory
    throw; // rethrow (dtor is not called, of course)
}
```

- of course, T::T () is assumed to be "safe": has no leaks

13

### On implementing strong guarantee

- sample of strong exception guarantee and roll-back

```
void doOperation (T const& someValue) {
    try {
        <update the state copying the giving value>;
        // e.g., the copy operation may fail & throw
    } catch (...) { // catch any exception
        <restore the old state and its invariants>;
        throw; // now rethrow the original
    }
}
```

- exception neutrality: T-related exceptions pass through
- strong guarantee may be very tricky or too costly to achieve; STL does not provide it for all its operations
- special C++ idioms support strong guarantee (see later)

14

### Example: exception safe constructor

```
Vector::Vector (size_t sz, T const& x) { // illustrative
    rep_ = (T*)::operator new (sz*sizeof(T)); // new may fail
    T * p = rep_; // element address
    try { // construct sz items
        for (; p != rep_ + sz; ++p) new (p) T (x); // ctor may fail
    } catch (...) { // handle T constructor failures
        while (p-- != rep_) // destroy all constructed items
            p->T::~~T (); // call T's destructor
        ::operator delete (rep_); // release memory
        throw; // propagate the original exception
    }
    size_ = capacity_ = sz; // OK: Vector initialized
}
```

15

### Exception safety (cont.)

- similar implementation for copy construction: if copy of an item fails, must destruct previously copied ones
- assignment operators can often be safely programmed with an existing copy constructor and *swap*

```
X& X::operator = (X const& rhs) {
    X tmp (rhs); // may fail
    swap (tmp); // does not fail
    return *this;
}
```

- here we trust that the *swap* operation does not throw
- the same requirement for X's destructor (*tmp* becomes destroyed at the end of the function before the return)

16

### Destructors are critical for exception handling

Question:

*What happens if an exception is thrown out from a destructor while the system is still propagating another?*

- destructors should not (generally) allow exceptions to escape from them
  - since propagating an exception calls destructors and if such a destructor lets its exception escape, the program is immediately terminated by system
- so a destructor should trap all local exceptions
  - handle and recover from the exception, or log out an error diagnostics and shut down the program
- the compiler cannot check, and so it is the programmer's responsibility

17

### Example: destructors "throwing exceptions"

```
struct B { // a class with bad behavior
    ~B () { // doesn't trap local exceptions
        bool b = std::uncaught_exception (); // for tracing
        throw string ("~B error"); // lets exception escape
    }
};

struct A { // has a "complicated" destructor that
    ~A () { // uses a B with its bad behavior
        try { B b; // ~B throws but
        } catch (std::string const& s) { // exception trapped
            assert (s == "~B error"); // "handled" internally
        }
    }
};
```

18

```

int main () {
    try { B b;                // ~B throws "~B error"
    } catch (string const& e) { // is caught
        assert (e == "~B error"); // matches ok
    } // throw out from ~B is handled ok
    try { A a;
        throw string ("A error"); // ~A handles ~B throw
    } catch (string const& e) { // original is caught
        assert (e == "A error"); // and matches ok
    } // local ~B exception inside ~A is handled OK
    try { B b;
        throw string ("error"); // ~B throws, too
        // => program stopped since ~B exception escapes
    } catch (...) { /* never comes here */ }
}

```

UNIVERSITY OF HELSINKI

19

### Case: how safety mechanisms work

- consider the following C++ class and code

```

class A : public B {
public:
    A () {} // implicit ctor calls
    X x; Y y; // two public members
}; ... // implicit dtor

A * a = new A; ...
delete a;

```

- the **A** constructor may seem empty but actually it handles the construction of **B**, **X**, and **Y** parts of an **A** object
- similarly **A**'s compiler-generated destructor handles the destruction of all these members

UNIVERSITY OF HELSINKI

20

```

struct A_Impl { // class A's hypothetical implementation
    // reserve space for the data members of A
    char b [sizeof (B)]; char x [sizeof (X)]; char y [sizeof (Y)];
    A_Impl () {
        new (&b) B; // call B::B () may throw
        try { new (&x) X; } // call X::X () may throw
        catch (...) {
            ((B*)&b)->~B (); throw; // destruct B part & rethrow
        }
        try { new (&y) Y; } // call Y::Y () may throw
        catch (...) {
            ((X*)&x)->~X (); ((B*)&b)->~B (); throw;
        } // otherwise: an A is now constructed OK
    }
    ~A_Impl () { // destruct all its members, in reverse order
        ((Y*)&y)->~Y (); ((X*)&x)->~X (); ((B*)&b)->~B (); }
}

```

UNIVERSITY OF HELSINKI

21

### Case (cont.)

```

A * p = new A; ... // create a dynamic A and use it
delete p; // later get rid of it

```

- using the class **A\_Impl**, above code is implemented as
- ```

void * p = ::operator new(sizeof (A_Impl)); // (1) allocate
// operator new throws std::bad_alloc upon failure
try {
    new (p) A_Impl; // (2) create an A at p (or fail)
} catch (...) {
    ::operator delete (p);
}
... // some other code
((A_Impl*)p)->~A_Impl (); // (1) release A resources
::operator delete (p); // (2) release p's space

```

UNIVERSITY OF HELSINKI

22

### Need for exception safety

Reusable library components vs. basic applications

- different levels of exception safety can be identified and are appropriate in different situations
- strong guarantee may be too expensive or not worth it: not all processing or programs can be made or need to be "failure safe"

For example

- an application program is not necessarily meant to be a separate reusable component (or a part of a library)
- when encountering an error, a simple application program may report errors, decide to end its execution, discard all calculated results, and require the user to try it again with more valid input

UNIVERSITY OF HELSINKI

23

|                           | Container-Operation Guarantees |         |                      |                              |
|---------------------------|--------------------------------|---------|----------------------|------------------------------|
|                           | vector                         | deque   | list                 | map                          |
| <i>clear()</i>            | nothrow                        | nothrow | nothrow              | nothrow                      |
| <i>erase()</i>            | (copy)                         | (copy)  | nothrow              | nothrow                      |
| <i>1-element insert()</i> | strong                         | strong  | strong               | strong                       |
| <i>N-element insert()</i> | (copy)                         | (copy)  | strong               | basic                        |
| <i>merge()</i>            | —                              | —       | nothrow (comparison) | —                            |
| <i>push_back()</i>        | strong                         | strong  | strong               | —                            |
| <i>push_front()</i>       | —                              | strong  | strong               | —                            |
| <i>pop_back()</i>         | nothrow                        | nothrow | nothrow              | —                            |
| <i>pop_front()</i>        | —                              | nothrow | nothrow              | —                            |
| <i>remove()</i>           | —                              | —       | nothrow (comparison) | —                            |
| <i>remove_if()</i>        | —                              | —       | nothrow (predicate)  | —                            |
| <i>reverse()</i>          | —                              | —       | nothrow              | —                            |
| <i>splice()</i>           | —                              | —       | nothrow              | —                            |
| <i>swap()</i>             | nothrow                        | nothrow | nothrow              | nothrow (copy-of-comparison) |
| <i>unique()</i>           | —                              | —       | nothrow (comparison) | —                            |

From Appendix E [Stroustrup]

UNIVERSITY OF HELSINKI

### Exceptions and ctors/dtors: summary

The following built-in C++ mechanisms enable resource management even in case of failures and exceptions

1. an exception throw causes the unwinding of call stack
  - all objects located between the places where the exception is thrown and caught are destroyed, i.e., their destructors are called
2. suppose that an exception is thrown inside a constructor, which has already constructed one or more members
  - the run-time system calls the destructors of the already constructed members to release resources reserved by those members
3. a failed "new X" operation always
  - releases the space allocated for the X object



25

### Summary

- write exception-safe class libraries and components
- three different levels of exception safety can be provided: *basic*, *strong*, and *nothrow*
- *exception neutrality* needed especially for templates (unknown type parameters with unknown exceptions)
- play safe to prevent bugs and to debug
  - make redundant checks to verify assumptions
  - always initialize everything (especially pointers) to minimize random and unpredictable states
  - remember to clean up resources
  - for raw pointers use *smart pointers* (discussed later)



26