

6 Architecture of C++ programs



Preview

- managing memory and other resources
 - "resource acquisition is initialization" (RAII)
 - using `std::auto_ptr` and other smart pointers
 - safe construction of an object revisited
 - on reference counting techniques
- physical structure of C++ programs
 - header files and physical dependencies
 - how to organize a C++ program into files
- on principles of object-oriented programming
 - programming to an interface
 - using Design Patterns
 - sample patterns: *Template Method*, *Singleton*



Unmanaged pointers

(1) *raw pointers* cannot have "destructors" that clean up

```
void f () { // case 1: uses a raw pointer
    X * ptr = new X;
    ... push_back (Student ("Joe")); ... // may throw
    // calling other functions that may throw or not . .
    // all code maintained and changing over time . .
    delete ptr; // possible never executed: potential leak
}
```

(2) ctor exception safety works only for *class-type* members

```
X::X () // case 2: uses pointer members p1 and p2
: p1 (new A), // may succeed or throw (but OK)
  p2 (new B) { // if this throws, p1 is never released!
} // note that p1 has no destructor to execute
```



UNIVERSITY OF HELSINKI

3

Unmanaged pointers (cont.)

(3) trying to recover from an exception

```
try { // case 3: use manual recovery
    typedef std::vector <int> IntStack;
    IntStack * stack = new IntStack; // local pointer
    stack->push_back (20); ... // may throw
} // .. potential leak
catch (std::exception const& e) {
    std::cout << e.what () << std::endl;
    // ... ?
}
```

- if *new* fails, no problem: space is released by the system
- but if *push_back* throws, reserved memory is here lost
 - move the pointer *stack* out of the *try* block (if can)
 - manual release is still a low-level and ad-hoc solution



UNIVERSITY OF HELSINKI

4

Resource management in C++

- to manage resources in C++, use the RAII principle
"Resource Acquisition Is Initialization"
- in general, your code may use several different resources, such as memory, files, locks, ports, etc.
- resources can have managers around them, implemented as classes ("resource wrappers")
 - resources are acquired by invoking a constructor that associates the resource with some data member
 - resources are released by destructors
- the manager provides operations to access and update the resource
- to prevent unwanted resource copies, the copy ctor and the assignment of the wrapper can be made *private*



UNIVERSITY OF HELSINKI

5

Resource management (cont.)

The use of resources, r_1, r_2, \dots, r_k , divided into three steps

1. acquire all resources r_1, r_2, \dots, r_k
2. actual use of resources
3. release resources r_k, \dots, r_2, r_1

Implemented in C++ as follows:

```
... Resource1 r1 ( ... );      // acquire a resource
    Resource2 r2 ( ... ); ...  // acquire another resource
    r1.useRes (); ...          // use these resources
} // automatic release of r2 and r1 (in reverse order)
```

- if, e.g. a ctor throws, all reserved resources are released
- works similarly for (1) declared local objects, (2) nested objects (data members), and (3) any dynamic objects *owned* by (1), (2), and (3) objects



UNIVERSITY OF HELSINKI

6

Smart pointers

Problem: built-in pointers may be uninitialized, no automatic release for raw C pointers, create memory leaks, point to destroyed objects, etc.

Solution: smart pointers provide a much better behavior

- can be used as wrappers for primitive pointers
- objects look and behave as if they were pointers
- guarantee initialization (zero or otherwise) & clean up
- an application of the RAII principle
 - always properly set up resources by constructors
 - destructors are automatically called when they go out of scope (or are deleted - or during exceptions)

You can overload two pointer operators

- operator `->`, for member access (call a method, etc.)
- operator `*`, for accessing the whole object (as an *lvalue*)



UNIVERSITY OF HELSINKI

7

Use of *auto_ptr*

- include the header file `<memory>` to get `std::auto_ptr`
- `new` allocates memory and creates an object
- initialize `auto_ptr` with the object created

```
Student * p = new Student (321);  
std::auto_ptr <Student> stud (p); // not: <Student*>
```

- better: always immediately pass the object to its owner

```
std::auto_ptr <Student> stud (new Student (321));
```
- `stud` becomes the *sole* owner of the object pointed to



UNIVERSITY OF HELSINKI

8

auto_ptr features

- requires that there should be *at most one* owner
- can be used and acts like a regular pointer

```
std::cout << stud->getNumber ();    // output "321"
```

- when it goes out of scope (or is otherwise destructed), the object it owns is destroyed

```
void fun () {  
    Student * ptr = new Student (783);    // bad idea  
    auto_ptr <Student> stud (new Student (321)); // ok  
    ...  
    std::cout << stud->getNumber () << std::endl;  
} // at end, deallocation for stud, but not for ptr
```



UNIVERSITY OF HELSINKI

9

auto_ptr features (cont.)

- the **reset** function resets the state of the wrapped pointer, by deleting the old object and assigning a new one

```
auto_ptr <Student> ps;    // initialized to zero  
ps.reset (new Student (783));    // ps owns an object  
ps.reset ();    // delete the object and set to zero (0)
```

- how to transform **auto_ptr** into a regular C pointer

```
T * get () const    // C-pointer to the owned object, or 0  
if (ps.get () == 0) ...    // ps does not own any object
```

- can also release the ownership of the object

```
auto_ptr <Student> ps (new Student (321)); ...  
Student * ps1 = ps.release ();    // make unmanaged  
// now used via regular pointer: no automatic deletion
```



UNIVERSITY OF HELSINKI

10

Memory management using *auto_ptr*

```
...
{
    typedef std::vector<int> IntStack;
    // use dynamic object as it were a local object
    std::auto_ptr<IntStack> s (new IntStack); ...
    s->push_back (20);           // may throw
    ...
    // no problem: automatic release at exception
}
```



UNIVERSITY OF HELSINKI

11

Safe construction using *auto_ptr*

```
class X {
public:
    X () : p1 (new A), p2 (new B) { } ...           // safe ctor
    X (X const& rhs)                                // safe copy ctor
        : p1 (new A (*rhs.p1)), p2 (new B (*rhs.p2)) { } ...
    X& operator = (X const& rhs) {                  // safe assignment
        X tmp (rhs);                                // may throw but OK
        p1 = tmp.p1; p2 = tmp.p2;                   // nofail guarantee
        return *this; }
    ...
private:
    std::auto_ptr<A> p1;                             // OK: auto_ptrs are
    std::auto_ptr<B> p2;                             // class-type data members
```



UNIVERSITY OF HELSINKI

12

auto_ptr features (cont.)

- the copy operators *transfers the ownership* of an object
`auto_ptr<Student> stud1 (new Student (321));`
`auto_ptr<Student> stud2 (stud1); // owned by stud2`
`auto_ptr<Student> stud3 = stud2; // owned by stud3`
- never can pass an `auto_ptr` object as a *value parameter*
`void ShowStudent (auto_ptr <Student> s) { .. // error`
`... // some nice display of s`
`auto_ptr <Student> stud (new Student (321));`
`ShowStudent (stud); // after this call, stud is 0`
`std::cout << stud->getNumber (); // error`
- unfortunately, the compiler may not complain..
- also, cannot use an `auto_ptr` with STL containers (they require conventional copy semantics for their elements)



UNIVERSITY OF HELSINKI

13

Reference counting

- uses a *handle-body* solution to share objects and deallocate them when they are no longer used by anyone
 - the handle keeps track of the number of references
 - have reference counters in the body - or as *separate shared* objects (also called "nonintrusive")
- define copying operations in the handle with the correct semantics (~ Java reference variables)
`a = b; // use reference semantics`
 - increment the count of the right hand side object
 - decrement and check the count of the left hand side
 - assign pointer to the left hand side object
- implemented using C++ templates, of course



UNIVERSITY OF HELSINKI

14

```

template <typename Body>           // simplified version
class CountedPtr {
public:
    explicit CountedPtr (Body *);    // bind owned object
    Body * operator -> ();           // to access members
    Body& operator * ();             // to access the whole object
    CountedPtr& swap (CountedPtr&);  // swap values
    // other operations . . .
    CountedPtr ();
    CountedPtr (CountedPtr const&);
    CountedPtr& operator = (CountedPtr const&);
    ~CountedPtr ();
private:
    Body * rep_;                    // the shared object
    std::size_t * count_;           // separate for generality
};

```



UNIVERSITY OF HELSINKI

15

Reference counting (cont.)

Apply reference counting to Student objects:

```

CountedPtr <Student> p1 (new Student (123));
std::cout << p1->getNumber () << std::endl;
CountedPtr <Student> p2 (p1);           // shares the same
CountedPtr <Student> p3 (new Student (321)); // new one
. . .
p3 = p2;                               // now shares the same
std::cout << *p2 << std::endl;         // print out object
std::cout << p3->getNumber () << std::endl; // print no
. . .
// the Student destroyed when the count becomes zero

```



UNIVERSITY OF HELSINKI

16

Reference counting (cont.)

```
template <typename Body>    // a sample implementation
CountedPtr <Body>::CountedPtr (Body * b)
: rep_(0), count_(0) {      // or INVPTR_VAL
    try {                   // prevent leak due to a throw from new
        rep_ = b;           // cannot fail
        count_ = new std::size_t (1); // may fail
    } catch (...) {         // or: std::bad_alloc const&
        delete rep_; rep_ = 0; // ownership has transferred
        throw;              // rethrow exception
    }
}
```



UNIVERSITY OF HELSINKI

17

Notes on the implementation

Why used zero-initialization of pointers above?

- it seems that: when a failure occurs, no object and no pointers are returned or left behind to misuse
- still, initializing primitive values reduces unpredictable values and behavior (less random pointers around)
- behavior depends on implementation
 - the address of an object can be assigned *before* its constructor is executed
`ptr = new CountedPtr . . // ptr is assigned before ctor`
`// and may be later "accidentally" used`
 - of course, the behavior is undefined and no absolute guarantees are achieved
 - some C++ programming environments do such initializations as a part of their debugging support



UNIVERSITY OF HELSINKI

18

Reference counting (cont.)

```
template <typename Body>    // a sample implementation
CountedPtr <Body>&          // uses reference semantics
CountedPtr <Body>::operator = (CountedPtr const& rhs) {
    if (rhs.count_)
        ++*rhs.count_;      // one more reference to rhs
    if (count_ && --*count_ == 0) { // check if lhs is garbage
        delete rep_; delete count_;
    }
    rep_ = rhs.rep_;         // share the object
    count_ = rhs.count_;     // share its counter
    return *this;
}
```



UNIVERSITY OF HELSINKI

19

Notes on the implementation

- can here omit the self-assign test (even if "this == &rhs"):
 // if (this == &rhs || rep_ == rhs.rep_) return *this;
 ++*rhs.count_; // increments rhs (/lhs) count_
 if (--*count_ == 0) { // decrements lhs (/rhs) count_
 – actually optimized by leaving out (mostly useless)
- some overhead for separately allocated counters
 size_t * count_; // shared reference count
- could be optimized by a special allocation strategy
 (overload *new* for a counter class [Stroustrup, p. 421])
- note that **auto_ptr** avoids any such overheads
- sometimes reference counts can be placed into the **Body**
 objects (say, for some *string* class implementation)



UNIVERSITY OF HELSINKI

20

Problem: unnecessary header dependencies

```
// File: data.h
class Data { ...
};

-----

// File: client.h
#include "data.h"    // gets the Data class definition
class Client {      // compiles OK, but bad style
    Data query () const; ...
private:
    Data * ptrData_;
};
```

- changes to **Data** propagate to **Client**-related source code
- the physical dependence may create maintenance problems - or, at least, force recompilations



UNIVERSITY OF HELSINKI

21

Required class information

What information is required from the class X in order to compile client code? For example:

```
X obj;    // compiler needs to know instance size
          // to allocate space for the object
```

However, this information *is not required* for:

- (1) members with pointers or references, e.g., **X *** or **X&**
- (2) function *declarations* with value parameters or results

```
X getX () or void print (X par) need only declaration
```

- only the *caller* of the operations needs the definitions to determine the required sizes (to actually pass values)
- thus, many times header files don't need to include full definitions to define their services and interfaces



UNIVERSITY OF HELSINKI

22

Breaking unwanted dependencies

```
// File client.h
class Data;                // forward declaration only
class Client {
public:
    Data query () const; ... // OK: no impl. needed here
private:
    Data * pData_; ...      // OK: no impl. needed here
};
```

- only source code that actually creates objects needs to include appropriate header files
 - e.g., "**client.cpp**" may need to include "**data.h**" (but no problem since it is an isolated translation unit)



UNIVERSITY OF HELSINKI

23

Idiom: Pimpl (Pointer to impl.)

- also known as the *Handle-Body* idiom
- the class definition, with its private and protected parts, is often unnecessarily used to compile the client's code
- leads to a physical coupling between use of abstraction and its implementation, making maintenance difficult
- to avoid recompilation of the client's code, separate interface and implementation, and include only those header files that really are necessary for the client

```
class Abstraction {          // the handle part
...
private:
    struct Impl * body_;     // hides the body
    ...                     // to be used in implementation unit
};
```



UNIVERSITY OF HELSINKI

24

Problems with templates and headers

- not all "classes" can be forwarded with names only
 - `std::string` is really a *typedef* of a template instance and thus cannot be introduced by its name only
- the standard library provides the header `<iosfwd>` with minimal declarations for stream templates and their standard *typedefs*, such as `std::ostream`
- similar practice is recommended for user-defined headers
- no such forward header file exists for `std::string`
- C++ implementations sometimes `#include` extra header files along system headers, making code nonportable
 - in another platform missing headers break compilation
 - minimize dependencies on nested `#includes` by always including system headers as the last ones



UNIVERSITY OF HELSINKI

25

Headers: summary

- `#include` only header files that are minimally needed to make your declarations and code to compile
 - let *users* include those header files *they need*
- prefer introducing classes by forward declarations (`"class X;"`)
 - sufficient for declaring any functions, and for using pointers or references as data members
- always first *#include* user-defined (custom) header files, and after them system header files
 - the strategy ensures that custom headers are properly defined and stay independent from any hidden extra header file inclusions



UNIVERSITY OF HELSINKI

26

OOP: programming to an interface

The basic principle of object-oriented programming:

Program to an interface, not an implementation.

- access objects in terms of their abstract interfaces
- the code looks something like:

```
Interface * p = getObject ();  
p->op (); ... // avoid binding to specific impl.
```

- the actual type of an object is not known (possibly not even yet defined)
- the code is made independent from implementations
 - but remember that raw pointers are problematic:
use smart pointers, such as `boost::shared_ptr` or the similar addition to the 200x version of C++ (C++0x)




UNIVERSITY OF HELSINKI

27

Inheritance vs. delegation

- reuse through inheritance (~ white-box reuse)
 - uses a compile-time relationship between two classes (base and derived); define extensions and override
 - cannot easily make modifications at run-time
- reuse through delegation (~ black-box reuse)
 - uses a run-time relationship between two objects



UNIVERSITY OF HELSINKI

28

Inheritance vs. delegation (cont.)

Object composition combined with inheritance is often a better alternative than just inheritance:

- the client uses an interface provided by the handle that delegates requests to the body
- the handle need not be changed for new implementations
- can define new functionality without affecting the client
 - when changes occur, the client's code need not be recompiled but only has to be relinked with the modified code



UNIVERSITY OF HELSINKI

29

Inheritance vs. delegation (cont.)

- delegation decomposes an abstraction into separately manageable parts
- the body may be polymorphic (implemented by derived classes) and replaced at run time
- several design patterns are based on delegation techniques: *Bridge*, *Proxy*, *State*, etc.
- in C++, handle objects are also used for modularization and for management of memory, consider
 - the *Pimpl* idiom
 - smart pointers that are defined for resource management (see the part on RAII)



UNIVERSITY OF HELSINKI

30

Source of Design Patterns

In software projects, we have concrete problems, e.g.:

- how to sequentially *access the elements* of an aggregate *without* the knowledge of its implementation
- how to represent a *shared* global *resource* as an object
- a way to provide for the *undoing* of actions
- how to provide a unified *interface to a set of services*

Design patterns are solutions to such common problems

- (1) the *name* of the pattern,
- (2) the *problem and its context* (circumstances, forces)
- (3) the *solution*, usually as a set of collaborating objects
- (4) the *consequences* (pros and cons) of using the pattern

Plus: sample code, implementation tips, related patterns, actual systems ("at least two"), etc.



UNIVERSITY OF HELSINKI

31

Source of Design Patterns (cont.)

- solutions to design problems that you see over and over
- design patterns are not invented but found (in systems)
- patterns give solutions to (nontrivial) problems; these solutions have been tried and approved by experts, who have also named, analyzed, and catalogued these patterns
- often provide a level of indirection that keeps classes from having to know about each other's internals
- give ways to make some structures or behavior modifiable or replaceable
 - usually, by objectifying some aspect of the system
- generally, help you write more reusable programs



UNIVERSITY OF HELSINKI

32

Philosophy of Design Patterns (cont.)

Different kinds of practices and reusable designs

- (1) **idioms** - describe techniques for expressing low-level, mostly language-dependent ideas (ref. counts in C++)
- (2) **design patterns** - medium-scale, mostly language-independent abstractions (solution "ideas")
 - usually depend on object-oriented mechanisms
 - essential features can often be described with a simple UML class diagram
- (3) **software frameworks** - consist of source code with variant parts, into which the user can plug-in specific code to provide the required structure and behavior
 - for example, GUI libraries in Java (with their Hollywood principle: use callbacks)



UNIVERSITY OF HELSINKI

33

Basic Design Patterns

- the **Template Method** makes part(s) of an algorithm changeable: define the skeleton of an algorithm, deferring some of its steps to subclasses (the related pattern **Strategy** uses delegation at run time)
- the **Iterator** pattern gives access to elements but hides internal organization (originally, a feature in **CLU**)
- the **Singleton** ensures a class only has one instance, and provides a global point of access to it (manages globals)
- the **Bridge** pattern separates interface and implementation hierarchies, and allows them vary independently (related to the **Handle-Body** idiom)
- the **Abstract Factory** provides an interface for creating families of related or dependent objects without specifying their concrete classes; uses **Factory Methods** or **Prototypes** to create the actual instances; the factory object is often a **Singleton**



UNIVERSITY OF HELSINKI

34

Template Method pattern


- the **Template Method** design pattern defines a general algorithm, in terms of some calculation steps that are expressed as abstract operations, to be defined in derived classes
 - do not confuse with generics and C++ templates
- the most basic pattern: almost trivial use of OOP
 - but still a *pattern* that identifies a *problem* to solve
- **Template Method** is a behavioral class pattern
 - two kinds of behavioral patterns
 - class patterns that use inheritance, and
 - object patterns that use object composition and delegation (e.g., **Bridge**, **Strategy**)



UNIVERSITY OF HELSINKI

35

Template Method pattern (cont.)



The **Template Method** design pattern

- abstract operations represent the variable parts
- template methods use these abstract operations
- e.g., an abstract operation could be a **Factory**
- the derived classes are responsible for implementing the abstract operations



UNIVERSITY OF HELSINKI

36

Template Method pattern (cont.)

- the template method itself usually should not be overridden, and therefore it is not defined as *virtual*
- all deferred operations that are used by the template method are defined as *virtual*, as well as *private* (or *protected*), since the client is not supposed to directly call them

Note *white-box reuse* strategy

- must know what needs to be redefined and what not



UNIVERSITY OF HELSINKI

37

The *Singleton* pattern

- used to constrain object instantiation to ensure that the client cannot create more than one object of a particular class
- in Java, the class `java.lang.Runtime` is a singleton class with a static method `getRuntime ()`, responsible for returning the single instance of this class
- an example of a *creational object pattern*
- in practice, is used to ensure that there is exactly
 - one window manager or print spooler,
 - a single-point entry to a database,
 - a single telephone line for an active modem, etc.



UNIVERSITY OF HELSINKI

38

```

class Singleton {
public:
    static Singleton& instance ();
    void op ();           // anything useful for the singleton
private:
    Singleton ();         // implement but make private
    Singleton (Singleton&); // don't implement but declare
    void operator = (Singleton&); // to avoid defaults
};
... // implementation file
Singleton * onlyInstance_; // = 0 by default
Singleton& Singleton::instance () {
    if (onlyInstance_ == 0)
        onlyInstance_ = new Singleton;
    return *onlyInstance_;
}

```



Singleton (cont.)

- returns a reference: harder to destroy accidentally
- as an alternative solution, can use *static local* variables

```

class Singleton {
public:
    static Singleton& instance ();           // reference
    ...
}; ...
// implementation file:
Singleton& Singleton::instance () {
    static Singleton instance; // created at first call
    return instance;
}                                     // and destructed after main

```



Notes on *Singleton* implementation

1. The client accesses features only through `instance ()`
`Singleton::instance ().op ();`
2. The static method `instance ()` uses *lazy evaluation*.
3. Public operations are application-specific services.
4. The constructors and the copy operations are defined as *private* or *protected*
`Singleton * s = new Singleton ();` // error
5. Often no use to delete and *recycle* a singleton (only one);
and the following would create a dangling pointer
`delete &Singleton::instance ();` // error
Solution: declare the destructor as *private*, too.
6. A *static* local object (within the function) is destructed after `main`; for the pointer solution, an automatic clean-up must be arranged separately (see next slide).



UNIVERSITY OF HELSINKI

41

```
static Singleton * onlyInstance_;    // = 0, by default
Singleton& Singleton::instance () {
    if (onlyInstance_ == 0)
        onlyInstance_ = new Singleton;
    return *onlyInstance_;
}
// deletion of a singleton for the pointer solution
struct SingletonCleanUp {            // helper class
    ~SingletonCleanUp () {
        delete onlyInstance_;
        onlyInstance_ = INVPTR_VAL;    // or: = 0
    }
};
static SingletonCleanUp Dtor_; // cleanup after main
```

- the singleton is destructed after the `main` terminates



UNIVERSITY OF HELSINKI

42

Application architecture: summary

- RAI: use objects to manage resources
- keep your headers clean and minimal
 - (1) don't reveal unnecessary implementation details
 - (2) use *class forward* declarations when sufficient
 - (3) but use always full namespace paths
- include only headers that really are needed to compile
- first include user-defined (custom) header files
- use namespaces, and *using* declarations (not *directives*)
- use *unnamed namespace* within implementation files
- use design patterns to make some aspects (structure or behavior) of software separate and manageable
 - *Singleton, Iterator, Template method, Factory*



UNIVERSITY OF HELSINKI

43