

# Optimization

Some slides courtesy of: M. Davies, A. Storkey, S. Boyd, P. Patrinos

# Optimization

- Why do we need it?
  - Demo:  
<http://www.benfrederickson.com/numerical-optimization/>
- How common it is?
  - <http://online.stanford.edu/course/convex-optimization-winter-2014>
  - MSc in optimization (Univ. of Edinburgh)

# Optimization in Learning

Most learning algorithms involve optimization, e.g.

- Minimizing an error function (Multi-Layer Perceptrons)
- Maximizing a Likelihood (Nonlinear regression)
- Minimizing an expected loss function (Bayesian Decision Th'y)

Can I use direct optimization? → Only for analytically solvable forms (e.g. quadratic)

In general analytical solutions **unavailable**.

Hence we need to use iterative schemes for optimization. e.g.

- Back-propagation (i.e. gradient descent)

# Iterative optimization

We will now look into the field of numerical optimization in more detail:

*Problem: minimize the function  $E(\mathbf{w})$  with respect to  $\mathbf{w}$*

The basis blocks of an iterative process in optimization:

- Currently at some position  $w_t$ .
- Choose a new position to go to  $w_{t+1}$ .
- Think of as choosing a direction to go in, and then a distance to go in that direction. Want to get as close to optimum.
- If we know the direction, but not sure how far, we can use line search.


# Maxima and minima

A – D : stationary points of

$$E(w) \quad (i.e. \quad \frac{\partial E}{\partial w} = 0)$$

D, B : minima

B : global minimum



A cost function may have multiple minima and other stationary points (including saddles). Recall minimum implies

1. gradient  $\nabla E = 0$

2. curvature  $\nabla^2 E$  is positive ( $v^T \nabla^2 E v > 0$ , for any  $v$ ).


In practice we generally have to make do with finding a local minimum.

# 1-D Optimization

## 1-dimensional versus multi-dimensional

1-D and multi-D optimization are fundamentally different.

Let  $x < y < z$  be three points on a 1-D cost function such that  $E(y) < E(x) \& E(z)$ .  
It follows that there must exist a minimum within the bracket  $(x, z)$ .



No equivalent point-wise property exists in dimensions greater than one!

# Bracket methods

Once we have a bracketed minimum we can proceed to find a new smaller bracket. There are various methods, e.g.

- **Golden section search:**

Start with the triple  $(a,b,c)$  then suppose we choose  $x \in (b,c)$  and evaluate  $E(x)$

- if  $f(b) < f(x) \rightarrow$  new triple  $(a,b,x)$
- otherwise new triple  $(b,x,c)$ .

How to choose  $x$  (and from which interval)? New bracket length will either be:

$$c - b \text{ or } x - a$$

Hence choose  $x$  in largest interval.

$$\text{either } x = a + \frac{3-\sqrt{5}}{2}(b-a) \text{ or } x = c + \frac{3-\sqrt{5}}{2}(b-c) \quad (\text{Golden mean})$$

convergence is linear (bracket guaranteed to be  $< 0.618$  smaller)

# Multi-dimensional methods:

## Local approximations


We can characterise a minimum locally by a quadratic approximation:

$$E(w_0 + \Delta w) \approx E(w_0) + \nabla E(w_0)^T \Delta w + \frac{1}{2} \Delta w^T \nabla^2 E(w_0) \Delta w + \dots$$

The gradient for  $E(w)$  can similarly be expressed as:

$$\nabla E(w_0 + \Delta w) \approx \nabla E(w_0) + \nabla^2 E(w_0) \Delta w + \dots$$

Suppose that  $w_0$  is a minimum, by expanding around it:

$$E(w_0 + \Delta w) \approx E(w_0) + 0 + \frac{1}{2} \Delta w^T \nabla^2 E(w_0) \Delta w + \dots$$


no gradient

Hence at a minimum  $\nabla^2 E(w_0)$  is positive definite ( $v^T \nabla^2 E(w_0) v > 0, \forall v$ ).

→ any perturbation *must* increase  $E(w)$ .



# Multi-dimensional methods: complexity and function evaluation

We will be looking at a number of different optimization techniques. We note the following costs:

$$\text{Evaluation of } E(w) \rightarrow O(NW)$$

$$\text{Evaluation of } \nabla E(w) \rightarrow O(NW)$$

$$\text{Evaluation of } \nabla^2 E(w) \rightarrow O(NW^2)$$

$$\text{Inversion of } \nabla^2 E(w) \rightarrow O(W^3)$$

Where  $N$  is the number of data observations and  $W$  is the number of weights.

# Multi-dimensional methods:

## Gradient descent with fixed step size

Recall gradient descent only uses linear approximation:

$$E(w_0 + \Delta w) \approx E(w_0) + \nabla E(w_0)^T \Delta w$$


For a fixed  $|\Delta w|$  best approximate reduction in  $E(w)$  is:

$$w^{(k)} = w^{(k-1)} - \eta \nabla E(w^{(k-1)})$$

where  $\eta$  is a (small) step size.

**Problem:** no concept of a minimum.

Is there necessarily a good single value for  $\eta$ ? **Answer:** NO!



# Multi-dimensional methods:

## Convergence in gradient descent

We first introduce a change in coordinates, using an **eigenvalue decomposition**

$$\nabla^2 E_0 = U \Lambda U^T, \Lambda - \text{diagonal}, U - \text{rotation}$$

So that  $v = U^T w$  and:

$$E \approx E_0 + [\nabla E_0^T U^T] \Delta v + \frac{1}{2} \Delta v^T [U^T \nabla^2 E_0 U] \Delta v + \dots$$

Now the contribution of each component of  $\Delta v$  is independent:

$$E = E_0 + \sum_k (b_k \Delta v_k + \frac{1}{2} \Lambda_{k,k} \Delta v_k^2)$$

(where  $b = \nabla E_0^T U^T = (U \nabla E_0)^T$ ).

For each component the optimal  $\eta = 1 / \Lambda_{k,k}$ , converging in a single step.

However gradient descent only allows us to set a single  $\eta$ .

# Multi-dimensional methods:


## Effect of step-size in gradient descent

- If we choose it too small convergence is *very* slow.
- If we choose it too large the algorithm may become unstable (oscillate around) and never converge.


If all the eigenvalues,  $\Lambda_{k,k}$  are similar then performance of gradient descent is okay.

If  $\Lambda_{1,1} / \Lambda_{W,W} \gg 1$ , convergence will be bad!

$\eta$  too small




$\eta$  too big



# Multi-dimensional methods:

## An example of gradient descent

- We start at  $[-0.1 \ 0.6]^T$



$$f(x_1, x_2) = 1 - e^{-(10x_1^2 + x_2^2)}.$$


# Multi-dimensional methods:

## Newton's method

Alternatively we can find the minimum of the quadratic approximation explicitly:

$$\Delta w = w^{(k+1)} - w^{(k)} = -\left[\nabla^2 E(w^{(k)})\right]^{-1} \nabla E(w^{(k)}) \implies w^{(k+1)} = w^{(k)} - \left[\nabla^2 E(w^{(k)})\right]^{-1} \nabla E(w^{(k)})$$

This is Newton's method (will need to iterate since quadratic is an approximation)



**Figure 9.16** The function  $f$  (shown solid) and its second-order approximation  $\hat{f}$  at  $x$  (dashed). The Newton step  $\Delta x_{nt}$  is what must be added to  $x$  to give the minimizer of  $\hat{f}$ .


# Multi-dimensional methods:

## Newton's method

Alternatively we can find the minimum of the quadratic approximation explicitly:

$$\Delta w = w^{(k+1)} - w^{(k)} = -\left[\nabla^2 E(w^{(k)})\right]^{-1} \nabla E(w^{(k)}) \implies w^{(k+1)} = w^{(k)} - \left[\nabla^2 E(w^{(k)})\right]^{-1} \nabla E(w^{(k)})$$

This is Newton's method (will need to iterate since quadratic is an approximation)





Gradient descent does not point directly towards minimum (in reality a stationary point), whereas the Newton direction does.

Both methods belong to a general iterative family with  $\Delta w = -M^{-1} \nabla E_0$ .

For gradient descent,  $M = \eta^{-1} I$ , for Newton  $M = \nabla^2 E_0$

$$w^{(k)} = w^{(k-1)} - \eta \nabla E(w^{(k-1)})$$

# Multi-dimensional methods: Newton's method



Compare this with the gradient decent result (few slides back is shown larger). The gradient points on a point of high gradient but not necessarily to a local stationary point. Thus, Newton methods have faster convergence but...they have problems. See next slide.




# Multi-dimensional methods:

## Problems with Newton Algorithm

There are a number of problems associated with a direct application of Newton's method in nonlinear networks

1. Evaluating and inverting Hessian is **expensive** ( $Order(NW^2)$  and  $Order(W^3)$ )
2. Convergence issues:

a) Newton is a zero-finding algorithm. May find a **maximum**.



b) May go **unstable** - step size may take the  $w$  outside the range where the quadratic approximation is reasonable

# Multi-dimensional methods: regularizing the Newton method

To stabilize the Newton step we can restrict the region of search (and get good of both worlds [Newton and Gradient Descent]).

If  $M$  is symmetric +ve definite then  $\Delta w = -M^{-1}\nabla E_0$  will point downhill.

Consider: 
$$\Delta w^{(k+1)} = -\left[\nabla^2 E(w^{(k)}) + \gamma I\right]^{-1} \nabla E(w^{(k)})$$

If  $\gamma \rightarrow 0$  then  $M \rightarrow \nabla^2 E_0$  (Newton step)

If  $\gamma \rightarrow \infty$  then  $M \rightarrow (1/\gamma)I$  (gradient descent,  $\gamma = 1/\gamma$ )

$\gamma$  controls the size of the search region. We can choose  $\gamma$  adaptively.

*e.g.*

If  $E(w^{(k)}) < E(w^{(k-1)})$  then

$$\rightarrow \gamma = \gamma \div 10$$

else

$$\rightarrow w^{(k)} = w^{(k-1)} \text{ and } \gamma \stackrel{5-19}{=} \gamma \times 10$$


# Multi-dimensional methods:

## Levenberg Marquardt method

We now specifically consider the sum-of-squared errors cost function.

$$E(w) = \frac{1}{2} \sum_n (f(x_n, w) - y_n)^2 = \frac{1}{2} \sum_n e_n^2$$

This term might make Hessian not +ve definite so we ignore it



With derivatives:

$$\nabla E(w) = \frac{1}{2} \sum_n \frac{\partial e_n^2}{\partial w} = \sum_n \frac{\partial f(x_n, w)}{\partial w} e_n$$

and

$$\nabla^2 E(w) = \frac{\partial}{\partial w} \left( \sum_n \frac{\partial f(x_n, w)}{\partial w} e_n \right) = \sum_n \left( \frac{\partial f(x_n, w)}{\partial w} \frac{\partial f(x_n, w)}{\partial w}^T + e_n \frac{\partial^2 f(x_n, w)}{\partial w^2} \right)$$

Ignoring second term Gives:

$$w^{(k+1)} = w^{(k)} - \left[ \sum_n \frac{\partial f(x_n, w)}{\partial w} \frac{\partial f(x_n, w)}{\partial w}^T + \gamma I \right]^{-1} \sum_n \frac{\partial f(x_n, w)}{\partial w} e_n$$

Note  $\frac{\partial f(x_n, w)}{\partial w}$  is just gradient (e.g. calculated using *backpropagation*).

# Multiple dimensions and line searches

In gradient descent we progressed a small way down in one direction and then selected a new direction, but can we **instead also** decide how much to move down?

**Alternative approach** - continue along direction until a *line minimum* is found. We already know how to do this using bracketing methods!

**Line searching:**

1. Choose a direction  $d_1$
2. Perform line minimisation on  $E(w_1 + \lambda d_1)$  *e.g.* using a bracketing method
3. Select new search direction and repeat.

How do we select the new direction? *Gradient Descent?*

# Line searches and gradient descent

A naive approach to line searching is to search along the steepest direction:  
However each new search MUST be orthogonal to the last.

Why orthogonal?


line search minimum

$$\frac{\partial}{\partial \eta_k} E(w_{k+1}) = \nabla E(w_{k+1})^T \cdot \frac{\partial w_{k+1}}{\partial \eta_k} = \nabla E(w_{k+1})^T d_k$$

$$\rightarrow \nabla E(w_{k+1})^T d_k = 0$$

new search direction,  $d_{k+1} = -\nabla E(w_{k+1})^T$ .

Hence  $d_k, d_{k+1}$  orthogonal



We are re-searching directions previously minimised (zigzag downhill)!

# Conjugate directions I

Can we construct directions that preserve the previous minimization work?

Amazingly the answer is yes....

The idea of conjugate directions is to choose a direction  $w = w_k + \lambda d_k$  such that:


$$\nabla E(w_k + \lambda d_k)^T d_j = 0, \forall j < k$$

Which implies (with quadratic approx.)

$$(\nabla E(w_k) + \nabla^2 E(w_k) \lambda d_k)^T d_j = 0$$

And hence

$$d_k^T \nabla^2 E(w_k) d_j = 0$$



# Conjugate gradient algorithm I

Suppose we have a quadratic function:

$$E(w) = E_0 + b^T w + \frac{1}{2} w^T H w$$

Starting at  $w_i$  and searching in direction  $d_i$  the line minimum is:

$$w_{i+1} = w_i + \alpha_i d_i$$

denoting the gradient at  $w_i$  as  $g_i = \nabla E(w_i) = b + H w_i$  we can solve for  $\alpha_i$

$$g_{i+1}^T d_i = (b + H(w_i + \alpha_i d_i))^T d_i = g_i^T d_i + \alpha_i d_i^T H d_i = 0$$

Hence:

$$\alpha_i = -\frac{g_i^T d_i}{d_i^T H d_i}$$

# Conjugate gradient algorithm II

We now choose a  $d_{i+1}$  that is conjugate to  $d_i$  we will try a modified gradient:

$$d_{i+1} = -g_{i+1} + \beta_i d_i$$

for some  $\beta_i$ . Solving for conjugacy gives:

$$(-g_{i+1} + \beta_i d_i)^T H d_i = 0$$

Hence:

$$\beta_i = \frac{g_{i+1}^T H d_i}{d_i^T H d_i}$$

In fact this choice of direction is conjugate with all  $d_j, j < i$ .

Finally we can write:

$$\beta_i = \frac{g_{i+1}^T (\alpha_i H d_i)}{d_i^T (\alpha_i H d_i)} = \frac{g_{i+1}^T (g_{i+1} - g_i)}{d_i^T (g_{i+1} - g_i)}$$

since  $g_{j+1} - g_j = H(w_{j+1} - w_j) = \alpha_j H d_j$  (no need to use  $H$ )



# Conjugate gradient algorithm summary

0. Choose initial weight  $w_1$  and search direction  $d_1 = -\nabla E(w_1)$

... at step  $j$

1. Find line minimum for  $E(w_j + \alpha_j d_j)$ , setting  $w_{j+1} = w_j + \alpha_j d_j$

(if not at minimum)

2. Evaluate new gradient,  $g_{j+1} = \nabla E(w_{j+1})$

3. Calculate new search direction,  $d_{j+1} = -g_{j+1} + \beta_j d_j$ ,

using the formula: 
$$\beta_j = \frac{g_{j+1}^T (g_{j+1} - g_j)}{d_j^T (g_{j+1} - g_j)}$$

4. repeat from step 3 (or after  $W$  steps begin again with step 2)

Note: the algorithm may also need to be iterated many more times (*c.f.* Newton method).

The search directions may deteriorate, therefore it is sensible to re-start the algorithm

every  $W$  steps (other strategies also exist)